COLLECTIVE: RESEARCH PROPOSAL

Can We Build Artificial Intelligence from Programmable Living Units?
Bridging Biological Computation, Molecular Manufacturing, and Emergent Collective Intelligence
Version 0.1 | 12 November 2026 | Open interdisciplinary discussion

ABSTRACT
COLLECTIVE is an exploratory research proposal asking whether heterogeneous programmable biological or biohybrid units could develop learned cooperation and, under suitable conditions, emergent organically embodied artificial intelligence. Each hypothetical unit combines local computation, addressable biological information storage, sensing, communication, metabolism, and a molecular-production interface. The proposal separates feasible component-level research from unproven integration and from the much stronger claim of emergent intelligence. The initial research program prioritizes critical literature review, falsifiable digital models, quantitative feasibility bounds, and institutionally supervised component experiments only if warranted.

This document is a concept paper, not evidence that the architecture works, a wet-lab protocol, or a claim that intelligence will emerge.

1. THE CENTRAL QUESTIONS

ENGINEERING QUESTION
Which interfaces can present-day science provide to combine computation, molecular manufacturing, and learned cooperation in programmable biological units?

SCIENTIFIC QUESTION
Could the resulting network, if constructed, exhibit adaptive capabilities not reducible to a fixed repertoire of preprogrammed behaviors? Could this justify describing the system as organically embodied artificial intelligence?

COLLECTIVE does not equate computation with intelligence, cooperation with consciousness, or a novel behavior with spontaneous evolution. Intelligence must be defined operationally using generalization, adaptive problem solving, transfer to unseen tasks, and robustness under controlled comparisons. Consciousness is not an assumed outcome or necessary success criterion.

A useful negative result would establish that a key interface is physically infeasible, or that distributed learning provides no advantage over simpler baselines.

2. CONCEPTUAL UNIT ARCHITECTURE

The starting unit is a hypothetical cell-like or compartmentalized biohybrid module. Initial dimensions are deliberately unconstrained; a coin-sized module is a thought experiment, not a claim that a single natural cell can be enlarged without transport and homeostasis constraints.

A. Addressable DNA archive: DNA encodes indexed information and potentially executable instruction sequences. The processor knows an address scheme and can request selected segments. DNA is NOT assumed to offer electronic random-access speeds. Reading, indexing, latency, errors, durability, and molecular access remain open problems.

B. Processor and RAM: a defined boot routine, instruction-set representation, local working memory, and an interface to retrieve encoded routines. An FFT routine could in principle be

loaded and executed if encoding, readout, and decoding are actually engineered. DNA storage alone does not execute algorithms.

C. Sensor interfaces: optical, electrochemical, chemical, mechanical, pH, and other environmental sensing, with calibration, bandwidth, noise, and feedback estimates.

D. Communication: an initially supplied common serial bus/protocol supports heterogeneous units, remote calculation requests, and synchronized measurements. The system is not expected to invent basic communication from scratch.

E. Metabolism: energy, material precursors, transport, repair, and waste disposal are explicit requirements. Electrical power is not interchangeable with ATP or molecular feedstock.

F. Molecular production: a speculative processor-directed manufacturing interface, possibly more general than a natural ribosome. Sequence specification, synthesis, folding, assembly, quality control, and recycling are distinct challenges. This subsystem is not currently a standard electronic port.

3. COMPUTATION AND COOPERATION

Units begin with different sensing and processing capacities. A complex task may be decomposed into subtasks; one unit measures, another selects sampling conditions, another executes spectral analysis, and another validates a proposed interpretation. The bus transfers data, uncertainty, and requests. Local memory and connection strengths change with feedback.

An illustrative computation is signal processing: estimate sensor bandwidth and relevant frequency range; select a sampling interval; retrieve an encoded FFT implementation; perform or delegate calculation; validate against noise and independent measurements. Parallel processing must account for communication overhead, synchronization, and energy costs.

Network topology, specialization, centrality, and hierarchy are observable outcomes rather than mandatory initial organizational structures. Links may strengthen or weaken based on measured utility, within bounded connectivity and resource budgets.

A central methodological challenge is disentangling actual collective learning from ordinary distributed computing. Baselines must include fixed-rule agents, non-learning agents, centralized computation, randomized network topology, and resource-matched alternatives.

4. REWARD DESIGN AND FAILURE MODES

The proposed ecology gives individual units scores and the collective a shared task outcome. Candidate reward dimensions include correctness, time to a verified answer, computational work contributed, energy/material cost, and measurable harmful side effects.

Rewarding raw computation alone invites gaming: a unit can perform unnecessary operations to earn credit. Rewarding only speed or low resource use can discourage difficult but essential analysis. A research program should compare policies and use ablations rather than assume a correct formula. Contribution should be estimated against task improvement, redundant work, and counterfactual baselines where possible.

Other anticipated failures include free-riding, communication flooding, premature consensus, overconfident sensor fusion, dominance by highly connected units, reward hacking, and accumulation of misleading memory. A hierarchy may emerge, but network centrality alone does

not prove authority, intelligence, or beneficial specialization.

No reward system is currently validated. Defining incentive mechanisms is itself a major independent research problem.

5. EXAMPLE TASKS AND FEEDBACK

Chemical example: reactive oxygen species (ROS), not ordinary molecular oxygen O2. The network might identify a reactive species, measure local chemistry, evaluate candidate responses, and detect byproducts. If a measured product is harmful, the proposed response receives negative feedback and an independent safety layer stops production. This does not eliminate risks from delayed or undetected products. It is a conceptual feedback example, not a biochemical protocol or therapeutic proposal.

Non-chemical examples should precede any physical intervention: adaptive signal analysis, distributed estimation, mathematical problem decomposition, sensor calibration, and coordination under communication limits.

To avoid circular claims of intelligence, held-out tasks must be genuinely unseen; initial stored programs and reference knowledge must be audited. A solution assembled from a stored blueprint differs from inventing a new physical capability.

6. THE EMERGENCE HYPOTHESIS

The strongest speculative question is whether a collective, instructed only to solve problems, could discover useful intermediate goals and physical adaptations without explicit commands to move, build an appendage, add hardware, or replicate. None of these behaviors is to be included as a prescribed objective.

Nevertheless, physical innovation requires a real action space: sensing, actuation, material handling, fabrication possibilities, and testing. A system cannot discover a physical capability that its environment and interfaces make impossible. The absence of an explicit instruction does not prove novelty; a general optimizer can produce unexpected actions without general intelligence.

Possible tests include unseen environmental constraints, blinded evaluation of new artifacts, exhaustive logging of accessible prior blueprints, comparisons with non-learning optimizers, and tests of whether a proposed structure improves independently measured task performance.

Whether such abilities appear is an experimental outcome, not an assumption.

7. FEASIBILITY GAPS

DNA interface: address selection, targeted reading, read speed, random access, molecular degradation, error correction, and instruction encoding.

Electronic-biological coupling: sensor transduction, noise, drift, electrical isolation, compatibility with wet biochemical environments, and repair.

Molecular manufacturing: reliable precursor delivery, selective chemistry, folding and assembly, toxicity monitoring, and reproducibility.

Metabolic integration: ATP and other energy carriers, thermodynamic budgets, transport across compartments, waste removal, and sustained homeostasis.

Distributed learning: nonstationary reward, contribution attribution, asynchronous signaling, limited bandwidth, and reliable shared semantics.

Scale and safety: physical containment, unintended interactions, observability, independent shutdown, institutional oversight, and clear boundaries against uncontrolled replication.

These gaps should be assessed separately before any integrated feasibility claim.

8. RESEARCH PROGRAM AND MILESTONES

PHASE 0: Critical literature review and feasibility matrix. For each subsystem record demonstrated capability, proposed extension, quantitative gap, relevant literature, and failure criteria.

PHASE 1: Digital models of heterogeneous agents with bounded energy, compute, memory, and communication. Benchmark cooperation, learning, reward gaming, and topology changes against centralized and fixed-rule controls.

PHASE 2: Interface-focused theoretical and engineering studies. Quantify DNA retrieval bandwidth and latency; sensor readout and uncertainty; molecular production limits; energy and mass balances.

PHASE 3: Modular, contained laboratory validation only with qualified teams, relevant approvals, and independent safety assessment. Test interfaces separately before considering integration.

PHASE 4: Decide whether an integrated system is scientifically justified. Document negative findings. Only after robust integration should claims about emergence be investigated.

Success criteria should be measurable and preregistered where feasible. Avoid treating a simulation demonstration as evidence of biological implementability.

9. TEAM AND COLLABORATION

A first review group could include a synthetic biologist or biochemist, a bioelectronics/computer architecture researcher, and an artificial-life/distributed-learning researcher. Later work may require microbiology, molecular data storage, materials science, control theory, chemical engineering, bioethics, and safety expertise.

Suggested initial deliverables: a shared glossary; a feasibility matrix; a minimal digital simulator; a risk register; and a prioritized list of falsifiable questions. This approach does not require establishing a wet laboratory at the outset.

Open collaboration should welcome skeptical critiques, prior-art references, negative results, and comparisons to simpler alternatives. Project governance, credit, licensing, and experimental oversight should be decided with collaborators before substantive work.

GitHub: https://github.com/altugkuran/COLLECTIVE-organic-AI
Contributions: open an issue stating the subsystem, established evidence, principal obstacle, and a safe discriminating test.

10. SCIENTIFIC POSITION AND LIMITATIONS

COLLECTIVE is not a working organism, not a validated artificial intelligence, not a proposal to bypass existing biology, and not a demonstration of self-organizing life. The combination of DNA storage, a digital processor, molecular manufacturing, and autonomous

collective learning is substantially beyond current integrated demonstrations.

The concept draws broad motivation from research areas including DNA data storage, synthetic biological circuits, bioelectronic interfaces, molecular manufacturing, artificial life, collective intelligence, and embodied intelligence. Specific prior-art claims and a verified bibliography are future work; this draft deliberately does not invent citations.

The principal research value may be in identifying impossible assumptions, building useful hybrid interfaces, or discovering that simpler systems suffice. The long-term question remains open:

CAN PROGRAMMABLE LIVING UNITS, THROUGH LEARNED COOPERATION, GIVE RISE TO ORGANICALLY EMBODIED ARTIFICIAL INTELLIGENCE?

This question should be answered by evidence, not by design ambition.